What Did Ten Points of Mean Recall Buy? Separating Group-Level from Case-Level Gains in Scene Graph Generation

Anonymous CVPR submission

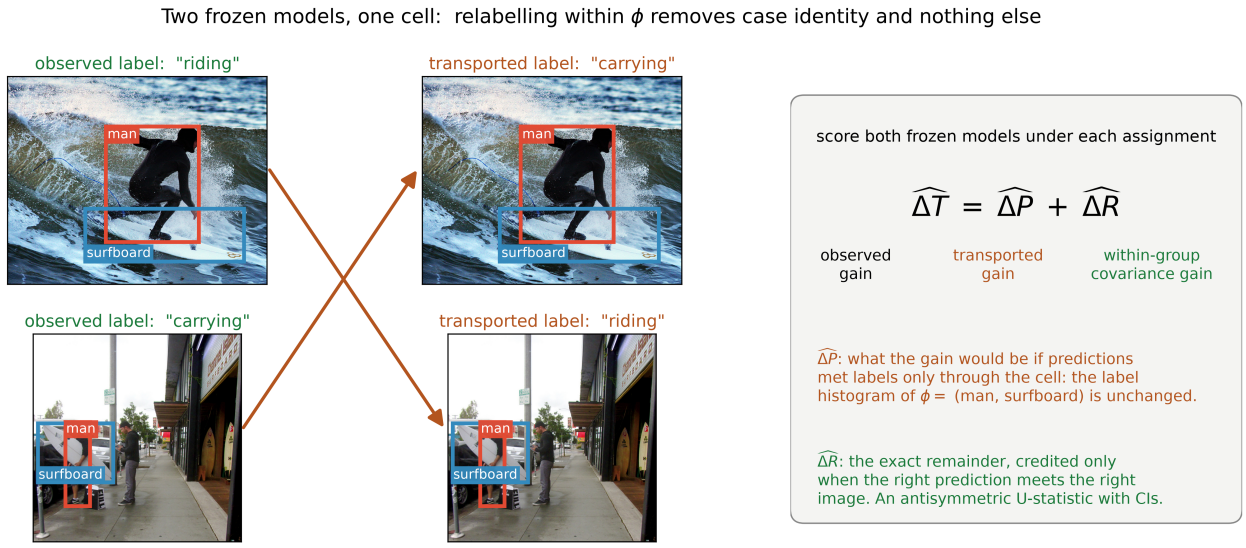


Figure 1. **Within-group relabelling.** Two Visual Genome relations share the subject–object pair (man, surfboard). Swapping their labels keeps the pair's predicate histogram exactly as it was but breaks the link between each image and its own label. Scoring two frozen models under the observed and the swapped labels splits their score gain exactly: $\widehat{\Delta P}$ is the part that survives the swap — earned at the level of the group — and $\widehat{\Delta R}$ is the covariance between each relation's prediction and its own label, which better ranking within the group can supply, as can a moved operating point.

Abstract

Debiasing methods for scene graph generation are judged by mean recall; Total Direct Effect (TDE) raises it in Visual Genome PredCls from 14.6 to 24.8. A model raises mean recall either by predicting rare predicates more often within each subject–object pair, or by telling which relations carry them, and the metric credits both alike. Scoring each prediction against the label of another relation from the same pair breaks the link between a relation and its own label but keeps the pair's label frequencies; this splits the difference between two frozen models, in a proper score or in mean recall, exactly into a group-level part and a within-pair covariance, with image-clustered intervals. We audit four released PredCls checkpoints, one run each. With class pairs as cells, 87% of TDE's mean-recall@50 gain (graph-constrained, micro-pooled; interval 83–91%) is group-level; the share is 84% for the official statistic and 18% with subject classes as cells. At $K = 50$ with the graph constraint a logit adjustment that changes no within-pair ranking recovers $+0.094$ of TDE's $+0.097$; at other cut-offs, or without the constraint, the two differ by 0.02 to 0.15. TDE's within-pair AUC shows no gain: four weightings exclude changes above 0.0038 to 0.028, tightly only under the two that favour head predicates. DLFE (a prior correction) and IETrans (relabelling and reweighting) gain $+0.111$ and $+0.220$, 94% and 97% group-level. PE-Net, a prototype-based model, gains $+0.170$, 86% group-level; its case-level part, $+0.024$, exceeds the $+0.001$ between two separately trained baselines, but its AUC change, $+0.005$, does not exceed their 0.006 spread, and without relation-level suppression its mean recall falls from 0.314 to 0.296.

1. Introduction

Scene graph generation is evaluated on a benchmark whose labels are unusually predictable from something every model sees. Given the ordered object pair (man, surfboard), the predicate on Visual Genome [21] is very often riding, and a lookup table of training frequencies is famously competitive [53]. Recall@K rewards models that echo this prior, so the field now leads with mean recall, which averages recall over predicates and so rewards getting the rare ones [5, 43]. Debiasing methods are now judged by how far they raise it. Total Direct Effect (TDE) [44], among the most cited of them, raises mean recall@50 of a MOTIFS model from 14.6 to 24.8 on the canonical VG150 split, in our evaluation of the released checkpoint.

What did those ten points buy? A model can move probability towards rare predicates *within each subject–object pair* — say carrying more often for (man, surfboard) — without becoming any better at deciding *which* relations are a man carrying rather than riding. Or it can learn to make that decision. The metric credits both alike. Existing remedies — head/body/tail-stratified recall [24], independent mean recall [26], zero-shot recall [29], rarity-stratified accuracy [20] — slice one model's performance differently; none attributes the gain *between* two models to one route or the other.

An exact split. Take two frozen models and a test set grouped by a discrete feature both observe, here the subject–object pair (Fig. 1). Score each model's prediction for one relation against the label of a *different* relation in the same group. This keeps each group's predicate frequencies, so gain earned at the level of the group survives it, and breaks the link between a relation and its own label, so gain earned on individual cases does not. The observed gain splits exactly into a *group-level* part $\widehat{\Delta P}$ and a *case-level* remainder $\widehat{\Delta R}$, a within-group covariance between the change in the predictions and the label (Theorem 1), for any proper score and for mean recall itself, with no fitted reference model, in $O(NK)$ on cached predictions, with image-clustered intervals.

What it finds. We re-ran TDE's released checkpoint, replayed the official evaluator exactly so that recall can be computed under any labelling, and split mean recall itself (Sec. 5). With class pairs as cells, 87% of the ten points is group-level. Because a thresholded metric also earns case-level credit when its operating point moves, we read the remainder against a control: subtracting the training predicate prior from the baseline's logits [31], known to rival debiasing [6, 46], changes no ranking within any group and recovers $+0.0941$ of the $+0.0972$ with the same split at $K = 50$ under the graph constraint (Tab. 2 gives the protocols where the two diverge). TDE is motivated as removing the harmful part of the training bias while keeping the good context [44]; on this checkpoint its gain is a change in which predicates are favoured within each pair, and DLFE (dynamic label frequency estimation) [6] and IETrans (internal and external data transfer) [54] show the same pattern.

Contributions. (i) An exact, model-free decomposition of the difference between two frozen models into a group-level part and a within-group covariance, the paired-gain form of Yates's covariance decomposition [50], for every Bregman score and for label-weighted hit rates such as mean recall, with cluster-robust inference, an operating-point control and a threshold-free within-group AUC (area under the ROC curve). (ii) Validation that the case-level part is zero when it should be and covers at the nominal rate under image clustering, with a map of what it does not separate: calibration, the operating point, the label mix and pair-varying prior shifts all move it, so ranking claims rest on the control and the AUC. (iii) Audits of three released prior-correcting checkpoints and one prototype-based checkpoint that split their mean-recall gains exactly, with the mechanism behind TDE's.

What is not claimed. The audits cover four released checkpoints, one run each, in predicate classification (PredCls: boxes and object labels are given, so every model sees the cell), with a TDE replication in scene graph classification and detection (SGCls, SGDet; Sec. 5): three prior corrections and one prototype-based model, nothing on one-stage or vision–language models. The split is conditional on the declared grouping, credits any signal varying inside a pair, shortcuts included, and moves with recalibration, the operating point and the label mix (Sec. 7); the official per-image statistic is split separately (Appendix Q.1).

2. Related work

Bias and its evaluation in scene graph generation. Since MOTIFS showed a frequency predictor of the predicate given the object pair to be competitive [48, 53], mean recall has become the headline metric, and debiasing methods span the causal TDE intervention [44], class-balanced sampling and group-wise learning [11, 13], predicate-rebalancing message passing [24], predicate-correlation and hierarchical losses [30, 49, 51], prototype embeddings [55], edge-selective and relation-aware transformers [19, 42], and correction or relabelling of the training labels [23, 54]. Correction by the predicate prior is known to rival them: post hoc, DLFE [6] recovers unbiased predicate proba-

bilities from estimated label frequencies and Structured Sparse R-CNN [46] applies logit adjustment [31]; in training, RTPB [4] resists a prior bias. The metrics themselves have been shown to reward pathologies [26], and open-vocabulary methods inherit them [25]. These works compare methods by their metrics; we ask which part of a metric's change is group-level. Outside scene graphs, VQA models collapse when answer priors change [1], models exploit non-visual shortcuts [15, 56], and gains are credited to mechanisms that did not produce them [39, 45]. Long-tailed recognition faces the same question for re-weighting and logit adjustment [3, 31]. All of these characterise a *single* model; we account for the gain *between* two.

Proper-score decompositions. Conditional on a grouping, a proper score splits into reliability and resolution terms [9, 32], for general proper scores [2]. Our case-level term is not that resolution term: it descends from the *covariance* decomposition of the Brier score due to Yates [50], which a monotone recalibration moves while the classical resolution is invariant to it. The two coincide only when both models are calibrated within every group. Variance estimators exist for single-model decompositions [40], not for a contrast between two models. The grouping loss [33] is the single-model relative of our case-level part: what calibration leaves behind because inputs sharing a confidence level differ; ours contrasts two models within groups declared in advance. Transporting labels within a group is a within-group permutation of the input–label pairing, so one model's case-level part is a conditional permutation importance [41] or model reliance [14] of the non-grouping part of its input, and the case-level part of a gain is the difference of the two models' (Proposition 2). What the split adds is the exact additive accounting of a paired gain, including thresholded metrics, and paired inference for it.

Comparing two predictors. At a single group our total-gain interval is the clustered Diebold–Mariano test [12, 16] (Corollary 3); the case-level term is a paired order-two U-statistic, as in DeLong *et al.*'s comparison of ROC areas [10, 18], with a Hájek projection [38] as its influence function.

3. Splitting a score gain

Setup. Let $q_0(\cdot \mid x)$ and $q_1(\cdot \mid x)$ be the predictive distributions of an old and a new frozen model over K labels, and $S(q, y)$ a bounded proper score, larger being better. Our default is the quadratic score $S(q, y) = 2q_y - \|q\|_2^2$, negative Brier loss up to a label-only constant; the log score is the second case we report. Let $C = \phi(X)$ be a discrete grouping both models observe, declared before the audit — for relation prediction, the ordered subject–object class pair. We call a value of C a *cell*. The gain vector

$$H_x(y) = S(q_1(\cdot \mid x), y) - S(q_0(\cdot \mid x), y), \qquad y = 1, \dots, K, \tag{1}$$

can be evaluated at a counterfactual label without rerunning either model, and that is all the construction needs.

Transport. For a cell c, let $\Delta T_c = \mathbb{E}[H_X(Y) \mid C = c]$ be the observed gain and $\Delta P_c = \mathbb{E}_{H|c}\mathbb{E}_{Y'|c}[H_X(Y')]$ the gain when each case is scored against an *independent* label Y' drawn from the same cell, and set $\Delta R_c = \Delta T_c - \Delta P_c$. Transport keeps the cell and its label distribution and removes only the pairing between a case and its own label.

Theorem 1 (Transport and covariance identities)**.** *For any integrable score contrast H,*

$$\Delta T_c = \Delta P_c + \Delta R_c,$$
$$\Delta R_c = \sum_{y=1}^{K} \mathrm{Cov}(H_X(y), \mathbb{1}\{Y = y\} \mid C = c)\,. \tag{2}$$

For the quadratic score, with $\Delta q_y(X) = q_1(y \mid X) - q_0(y \mid X)$,

$$\Delta R_c = 2 \sum_{y=1}^{K} \mathrm{Cov}(\Delta q_y(X), \mathbb{1}\{Y = y\} \mid C = c)\,. \tag{3}$$

Consequently, if the gain vector is a function of c alone, $\Delta R_c = 0$.

Eq. (3) is what the split means. The new model earns case-level credit only insofar as its probability *change* rises on the label the individual case actually carries, relative to other cases in the same cell. Anything constant within a cell — shifting probability towards rare predicates for every (man, surfboard) relation alike — lands in ΔP, which is why we call it the group-level part. It is *not* simply "prior fit": under the quadratic score ΔP_c is the improvement in how close the model's mean forecast sits to the cell's label frequencies *minus* the increase in the dispersion of its forecasts (Proposition 1). The covariance form is due to Yates [50] and generalises: for every Bregman score, including the log score, ΔR_c is the same covariance with Δq_y replaced by the contrast of the y-th gradient coordinate (Theorem 2). All proofs are in the supplementary.

Why it matters which part a gain is in. With $m_y(z) = \mathbb{E}[H_X(z) \mid Y = y, C = c]$, $\Delta R_c = \sum_{y<z} p(y)p(z)\, D(y, z)$ where $D(y, z) = m_y(y) - m_y(z) + m_z(z) - m_z(y)$: a label shift fixes every D and moves only the weights, whereas a group-level gain can turn into a loss, so a case-level gain is less exposed to a shift, though with three or more labels it can still reverse (Appendix C works a two-label example).

Estimator. In a cell with $n_c \geq 2$ cases, write $H_i(y) = H_{x_i}(y)$ and estimate

$$\widehat{\Delta T}_c = \tfrac{1}{n_c}\textstyle\sum_i H_i(y_i), \qquad \widehat{\Delta P}_c = \tfrac{1}{n_c(n_c-1)}\textstyle\sum_{i\neq j} H_i(y_j), \tag{4}$$

$$\widehat{\Delta R}_c = \tbinom{n_c}{2}^{-1}\textstyle\sum_{i<j} A_{ij}, \tag{5}$$

$$A_{ij} = \tfrac{1}{2}\{H_i(y_i) + H_j(y_j) - H_i(y_j) - H_j(y_i)\}. \tag{6}$$

These satisfy $\widehat{\Delta T}_c = \widehat{\Delta P}_c + \widehat{\Delta R}_c$ exactly in every sample, and $\widehat{\Delta R}_c$ is an unbiased order-two U-statistic [18] for ΔR_c (Theorem 3). Transporting only to the *other* cases' labels matters: the all-case plug-in shrinks the case-level gain by $(n_c - 1)/n_c$ (Proposition 5). With m_{cy} the count of label y in cell c, case i's transported score is

$$P_i = \frac{\sum_y m_{cy} H_i(y) - H_i(y_i)}{n_c - 1}, \tag{7}$$

so the whole estimator costs $O(NK)$ on cached predictions, with no fitted model, held-out fold or smoothing parameter. Cells are weighted by size; singletons cannot identify a within-cell counterfactual and are excluded and reported.

Inference. The Hájek influence contributions of the U-statistic (Appendix F), summed within each image, give one-way cluster-robust intervals, since relations from one photograph are not independent; the estimator is asymptotically normal under bounded scores and non-dominating clusters (Theorem 4), and at a single cell the total-gain interval is a clustered Diebold–Mariano test [12] (Corollary 3).

Confidence matching. Because ΔR is a covariance between probability *movements* and labels, rescaling a model's confidence moves score between the two parts while changing no ranking (Sec. 4), so models that differ in calibration are not compared on raw outputs. We split the images at random, fit each model a recalibration by maximum likelihood on one half, and split the gain on the other. The family is chosen by one criterion stated in advance of the comparison it serves (Sec. 5 dates the declaration): it must be rich enough that a model differing from another by a shift of each label's log-odds that is the same in every cell, which changes no ranking, has no case-level part under it. A temperature alone fails this, a temperature with a per-class bias meets it exactly and still preserves every within-cell ranking, so the latter is the primary protocol and the temperature-only split is reported beside it.

Splitting mean recall. Nothing in Theorem 1 or the estimator uses properness: both hold for any contrast $H_x(y)$ that can be evaluated at a counterfactual label. Mean recall@k is one. Let $h_x^{(m)}(y) \in \{0, 1\}$ say whether model m would recall relation x among its top k if its label were y; replaying the benchmark's matching rules on cached scores gives it for every y (Appendix O). Pooling relations, mean recall is $\sum_i w(y_i)h_{x_i}(y_i)$ with $w(y) = 1/(Kn_y)$, so $H_x(y) = Nw(y)[h_x^{(1)}(y) - h_x^{(0)}(y)]$ makes $\widehat{\Delta T}$ its change; transport preserves every n_y, and restricting w to one tier or predicate splits that share exactly. One difference matters for reading the result. A hit is a thresholded decision, so moving a model's operating point — subtracting a log-prior from its logits, say — flips first the cases it already scored as most like the favoured label, and those carry that label more often: a case-level part with no new discrimination. We therefore read the split against an *operating-point control*, the old model's logits shifted by the training log-prior [31], and pair it with a threshold-free test, the *within-cell AUC*: the probability that a case labelled y has larger log-odds $\log q(y) - \log q(z)$ than one labelled z in the same cell. Averaged with weight $n_c(y)n_c(z)$, the number of comparisons, it weights a cell by about n_c^2; with weight $n_c(y)n_c(z)/n_c$ it combines cells as the split does; we take the latter as primary and report both. A temperature, and any change constant within a cell, leaves either unchanged.

Choosing the grouping. The split is conditional on ϕ. Merging cells changes it by an explicit between-cell covariance term of either sign (Proposition 3), so the finest grouping the benchmark's construction defends is the conservative choice. A relevant variable left out of ϕ is the same term with the fully adjusted cell as the fine partition; Proposition 4 bounds it by one elicited correlation, after the robustness value of Cinelli and Hazlett [7].

4. Does the split measure what it claims?

A released checkpoint cannot validate the split, because what the model can see is unknown. We validate it where the answer is known, first in simulation and then on Visual Genome predictors whose information content we control. All intervals are 95%.

Simulation. With 24 cells over five labels and a new model interpolating from the cell prior towards an oracle, the case-level gain grows linearly with the injected signal, and under a within-cell null the randomization test rejects in 0.030 of 200 datasets at level 0.05. With relations clustered into images, the image-clustered interval covers in 94.0% of 500 runs (88.6% treating relations as independent); with VG150's own cell sizes it covers 96.3%, as do the mean-recall and share intervals, and only with two cases in every cell does it fall to 81.5% (Appendix J). A *prior-only* improvement gives a case-level gain centred at zero (0.0002, sd 0.0035). An *unrecorded within-cell shortcut* the new model uses is credited to the case-level part

(+0.041, sd 0.004): the split measures all within-cell signal, shortcut or not. A *calibration-only* change, a temperature-softened copy of an overconfident model, is credited a spurious case-level *loss* of −0.488 on raw outputs, which confidence matching reduces to −0.0003 (sd 0.0017). A fixed logit offset per label changes no within-cell ranking yet survives temperature matching (+0.0106) and vanishes under a temperature with a per-class bias (0.0000, sd 0.0003; Appendix Q.4): the criterion of Sec. 3 at work.

Controlled predictors on VG150. We train predicate classifiers whose inputs we control on a VG150-style split reconstructed from the released annotations (same vocabulary, images divided 70/30 between fitting and audit; Appendix N.1), not the canonical split of Sec. 5: a frequency table on the class pair (FREQ), a multilayer perceptron (MLP) on class embeddings (MLP-CLASS), and MLPs adding twelve box-geometry features (MLP-SPATIAL, MLP-SPATIAL+). The audit covers 227,337 of 229,605 test relations (99.0%) in 6,346 class-pair cells.

MLP-CLASS beats FREQ by 0.03386 of quadratic score, all group-level: the case-level gain is 0.00000 with a zero-width interval, as Theorem 1 forces for two functions of the class pair. Adding box geometry yields a case-level gain of 0.01439 (CI $[0.01373, 0.01504]$, $p = .002$) that no aggregate isolates; against MLP-CLASS the total gain of 0.00885 even *understates* it, because the group-level part falls by 0.00554.

5. What TDE's ten points bought

Tang *et al*. [44] release a causal MOTIFS PredCls checkpoint that yields two models from one set of weights: the baseline (MOTIFS-SUM), whose predicate logit sums a union-region visual term, a context term and a frequency prior indexed by the object pair, and TDE, which subtracts the same sum with an averaged context and so keeps only the context term minus its average. We ran it once with the authors' code on the canonical VG150 test split, saving every branch's logits, and replayed the official evaluator offline (Appendix O). The replay reproduces the official numbers to the fourth decimal: baseline $R@50 = 0.6612$ and $mR@50 = 0.1459$, TDE 0.4588 and 0.2476; without the graph constraint TDE's mean recall *falls*, from 0.3260 to 0.2981. The audit covers 183,642 relations over 26,446 images in 7,127 cells; 98.9% fall in the 5,044 cells holding at least two.

An operating-point control. We add one model that needs no training, LA: the baseline's logits minus $\tau \log \pi(y)$, the training predicate prior, with $\tau = 1$ [31]. Within a cell this adds the same constant to every log-odds $\log q(y) - \log q(z)$, so it changes no within-cell ranking of predicates: whatever case-level part it shows comes from a moved operating point, not from new discrimination. It reaches $mR@50 = 0.2343$ at $R@50 = 0.6437$. (It is not the inverse-frequency scaling of logits that barely moves SGG mean recall [28].)

Table 1. **Mean recall@50 split** (graph constraint, micro-pooled, class-pair cells) against the MOTIFS baseline, except TDE−LA. LA: the logit-adjusted baseline; DLFE: against its own released baseline; IETrans, PE-Net: released checkpoints against this baseline, separately trained; base_D: the baseline released with DLFE against this one, a floor for such comparisons. One released run each; intervals (image-clustered) cover test sampling only. ΔAUC: within-cell AUC change, weighted as the split weights cells (r, primary) or by comparison (c), 1,000 replicates; LA's is below 10^{-6}. Lower panel: TDE with other cells, and the official per-image statistic split by label permutations within cells.

	ΔmR	group	case [CI]	ΔAUC_r	ΔAUC_c
TDE	+.0972	+.0841	+.0130 [+.0093,+.0168]	−.0018 [−.0043,+.0016]	+.0025 [−.0028,+.0087]
LA	+.0941	+.0802	+.0139 [+.0101,+.0176]	0	0
TDE−LA	+.0031	+.0039	−.0008 [−.0050,+.0033]	−.0018 [−.0043,+.0016]	+.0025 [−.0028,+.0087]
DLFE	+.1112	+.1046	+.0065 [+.0010,+.0120]	+.0029 [+.0001,+.0055]	+.0018 [−.0042,+.0070]
IETrans	+.2204	+.2143	+.0060 [−.0031,+.0151]	−.0238 [−.0272,−.0187]	−.0093 [−.0185,−.0001]
PE-Net	+.1703	+.1460	+.0243 [+.0197,+.0289]	+.0050 [+.0008,+.0091]	+.0055 [−.0033,+.0149]
base_D	+.0206	+.0199	+.0007 [−.0016,+.0030]	−.0060 [−.0099,−.0011]	+.0017 [−.0072,+.0109]
TDE, other cells or statistic					
subject	+.0970	+.0175	+.0794 [+.0730,+.0859]		
+position	+.0973	+.0873	+.0100 [+.0065,+.0135]		
per-image	+.1024	+.0855	+.0168 [+.0111,+.0226]		

Protocol. Unless a number says otherwise: cells are subject–object class pairs; mean recall is $mR@50$ with the graph constraint, micro-pooled over identified relations; intervals are image-clustered; proper scores are matched by a temperature and a per-class bias per model, fitted on half the test images; the AUC is relation-weighted. These two primary quantities and the protocol were declared after the TDE and IETrans results were seen and before the DLFE outputs existed, so for TDE and IETrans they are a post-hoc choice and for DLFE and PE-Net a prior one; every other protocol, score and weighting is collected in Tab. 27, where, with that many configurations, one secondary interval excluding zero is not read as a finding.

Mean recall, split. Pooled over identified relations, mean recall@50 rises by 0.0972 (0.1017 on the official per-image average). Of that, +0.0841 is group-level and +0.0130 case-level, interval $[+0.0093, +0.0168]$

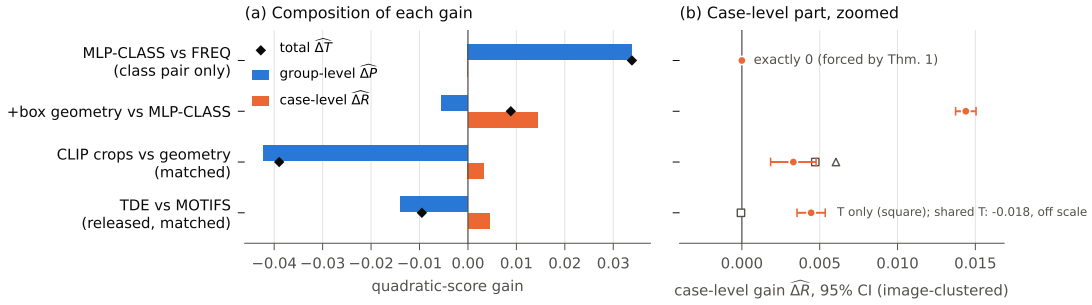


Figure 2. **Where each proper-score gain lives** (quadratic score). (a) Total gain $\widehat{\Delta T}$ (diamond) split exactly into group-level $\widehat{\Delta P}$ and case-level $\widehat{\Delta R}$. (b) $\widehat{\Delta R}$ with its image-clustered 95% interval. Class-pair-only models differ purely at group level (Theorem 1); geometry and CLIP crops buy case-level gain, CLIP despite losing overall; TDE's is group-level, with a small case-level residue. CLIP and TDE rows are matched by a temperature and a per-class bias per model (square: temperature only; triangle: one shared temperature).

([+0.0088, +0.0173] by an image jackknife; Tab. 1 and Appendix J): 87% of the gain, interval [83%, 91%], survives shuffling every label among the relations that share its object pair. The share is a property of the declared cells, not of TDE alone (Tab. 1, lower panel): refining the pair by the object's position or the subject's size leaves it at 90%, the official per-image statistic gives 84%, but cells of the subject class alone, which no longer hold the object's identity fixed, give 18% (Appendices K and Q.1). It lives in the body of the distribution: by training frequency (9 head predicates above 10,000 instances, 26 body, 15 tail below 500), the body gains +0.1224, of which +0.1082 is group-level, the head loses 0.0255, and the tail gains 0.0003 (the baseline recalls no tail predicate; TDE recalls the predicate to at 0.016). Figure 3 shows the mechanism one predicate at a time. Recall of parked on rises from 0.000 to 0.886, almost all of it group-level: TDE now says parked on for (car, street). Recall of on falls by 0.558 in the same way.

The case-level part is matched by the control. The control reproduces the whole picture: +0.0941 of mean recall, +0.0802 group-level and +0.0139 case-level. Against it, TDE's case-level part is −0.0008, interval [−0.0050, +0.0033]. When the boundary between on and parked on moves, the cases that cross it are those the baseline already scored as more parked on-like, and they are more often labelled so: the covariance is real, but bought with discrimination the baseline had. The direct test is the within-cell AUC, which no prior shift or recalibration can move: the control's change is below 10^{-6}. Weighted as the split weights cells, the primary weighting, TDE's change is −0.0018, [−0.0043, +0.0016]; by comparison +0.0025, [−0.0028, +0.0087]. Both put most weight on head predicates (57% and 65% on head–head pairs), so we add two that do not: every predicate weighted equally, as mean recall weights them, and every predicate pair weighted equally (Appendix Q.4). With 1,000 image-bootstrap replicates the 90% intervals exclude a change larger than 0.0038 under the primary weighting, 0.0077 by comparison, 0.023 per pair and 0.028 per predicate. Box geometry buys 0.063, 0.032, 0.138 and 0.069 over the class pair under the four (Sec. 6); the bounds pass a 25% margin of that under the first three and a 10% one under the primary weighting only, and the per-predicate weighting, the closest to mean recall, passes none: tail–tail comparisons are too few to decide (−0.097, [−0.228, +0.024]). The two narrow bounds rest on head-heavy weightings, and the margins were set after TDE's results were seen. No weighting shows TDE gaining discrimination (the baseline's AUC, 0.61 and 0.57, left room); per pair it falls, −0.0153, [−0.0249, −0.0004],

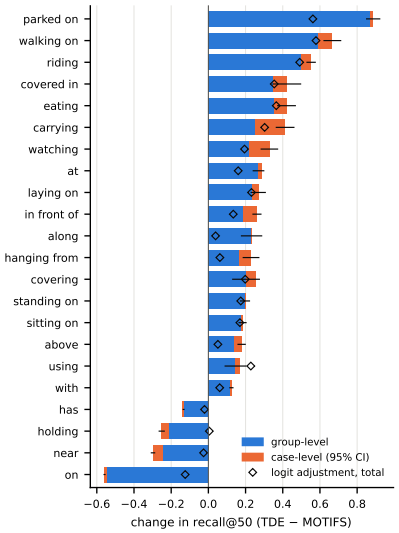


Figure 3. **Where TDE's mean recall went**, for every predicate whose recall@50 moved by at least 0.12: group-level part (blue), case-level part (orange) with its image-clustered 95% interval, and the total under logit adjustment of the baseline (diamond), which changes no within-cell ranking.

an exploratory result smaller than the 0.018 by which two separately trained baselines differ under that weighting. Against the control the mean-recall case-level part excludes only changes above 0.0043, a third of TDE's own, too wide to call it zero.

The proper score, matched. Under the primary protocol (Sec. 3), which makes the control identical to the baseline, TDE's quadratic-score change is -0.00951: -0.01396 group-level and $+0.00445$, $[+0.00354, +0.00536]$, $p = .001$, case-level (Fig. 2; cross-fitted $+0.00399$). That part is not a ranking measure: removing TDE's frequency prior, which leaves every within-cell ranking as it was (AUC change -0.00001), gives $+0.00817$, $[+0.00757, +0.00877]$, under the same protocol, so ranking is read from the AUC alone. The protocol also decides the sign: with a temperature alone the part is -0.00006, $[-0.00120, +0.00109]$ ($p = .555$), and with one temperature shared by both models -0.0184 (Appendices Q.1, Q.2 and Q.4).

What that licenses. On this checkpoint the ten points of mean recall are a change in which predicates the model favours within each object pair. A one-line prior shift that ranks the relations of every cell exactly as the baseline did reproduces the gain, its split and its tier profile at under a tenth of TDE's cost in $R@50$ (0.0175 against 0.2023). Table 2 gives all six recall protocols: the share runs from 77% at $K = 20$ to 86% at $K = 100$ under the graph constraint; without it TDE's mean recall falls, and its case-level part over the control turns positive at $K = 100$, $+0.0148$, $[+0.0064, +0.0232]$. The control's match is specific to $K = 50$ under the constraint: TDE leads it by $+0.0300$ at $K = 100$ and trails it by 0.0234 at $K = 20$ and by 0.08–0.15 without the constraint, where the total gain is negative and we report parts only.

How to read the split. The AUC audits the order of predicates for a labelled pair; a hit also depends on how the pair ranks among an image's pairs and on the mass given to *no relation* (0.855 for the baseline, 0.021 for TDE, 0.651 for the control), so the case-level part of mean recall can move with those too. Group-level is not a lesser gain: retrieval by predicate wants rare predicates named more often for a pair, captioning wants the right predicate on the right relation. The case-level part also credits shortcuts, operating-point movement and label noise, so discrimination claims rest on the control and the AUC; an AUC gain beyond the bounds above would change TDE's verdict. Nor is the share forced by the metric: an oracle gains $+0.5068$ in quadratic score, all case-level ($+0.8041$ against -0.2973). And it depends on the label mix: along a path to uniform predicate frequencies within each pair, TDE's temperature-matched quadratic case-level part rises from -0.00006 to $+0.01430$ while its AUC change stays undetected (Sec. 6).

Table 2. **The split under all six recall protocols**, micro-pooled mean recall with (mR) and without (ng-mR) the graph constraint at $K = 20, 50, 100$; case-level intervals image-clustered.

	TDE vs baseline			TDE vs LA	
Protocol	ΔmR	group	case [CI]	ΔmR	case [CI]
mR@20	+.0598	+.0462	+.0135 [+.0101,+.0170]	−.0234	+.0013 [−.0024,+.0050]
mR@50	+.0972	+.0841	+.0130 [+.0093,+.0168]	+.0031	−.0008 [−.0050,+.0033]
mR@100	+.1280	+.1099	+.0182 [+.0144,+.0220]	+.0300	+.0034 [−.0010,+.0078]
ng-mR@20	−.0037	−.0134	+.0096 [+.0058,+.0135]	−.0792	+.0013 [−.0035,+.0060]
ng-mR@50	−.0337	−.0464	+.0127 [+.0080,+.0173]	−.1272	+.0072 [−.0001,+.0145]
ng-mR@100	−.0489	−.0724	+.0236 [+.0175,+.0296]	−.1549	+.0148 [+.0064,+.0232]

Why. TDE's change is three subtractions whose splits add up exactly (Appendix O.1); the averaged context supplies $+0.0817$ of the 0.0972, and that term is, to 99.96% of its energy, one vector shared by every relation, correlated $+0.75$ with the log training prior: in PredCls TDE is to first order a learned logit adjustment. On zero-shot recall, where TDE beats the control by $+0.035$, its case-level part, $+0.013$, comes with an AUC that *falls* on comparisons involving an unseen triplet, -0.015, $[-0.020, -0.006]$ (Appendix Q.7).

Two more methods. DLFE [6] divides predicate probabilities by estimated label frequencies, a per-predicate log-odds shift, so a high group-level share is expected, and is released with a baseline trained alongside it (Appendix Q.5). Its $mR@50$ gain, $+0.1112$, is $+0.1046$ group-level and $+0.0065$, $[+0.0010, +0.0120]$, case-level. The correction itself changes no ranking: against the same network before it the AUC change is zero and the case-level part $+0.0054$, $[-0.0001, +0.0109]$, so the $+0.0029$ AUC gain over the release baseline, $[+0.0001, +0.0055]$, is a training difference. IETrans [54] relabels training data from general to specific predicates within an object pair and adds the training log-prior at test time; its released Neural-Motifs PredCls checkpoint is the IETrans+Rwt variant, also trained with a reweighted loss, re-scored the same way (Appendix O.2). It reaches $mR@50 = 0.3587$: $+0.2204$ over the baseline, $+0.2143$ group-level and $+0.0060$, $[-0.0031, +0.0151]$, case-level. Its relation-weighted AUC falls by -0.0238, $[-0.0272, -0.0187]$, and by comparison barely, -0.0093, $[-0.0185, -0.0001]$. A fall measures agreement with the recorded labels, and IETrans relabels general predicates as specific ones, so part of it may be synonymy rather than lost discrimination. No IETrans baseline was released: this compares two separately trained models and does not isolate the effect of relabelling; merging synonymous predicates leaves the loss in place (Appendix Q.8).

A method that is not a prior correction. PE-Net [55] scores a relation by its similarity to learned predicate

prototypes, with no prior. Its released checkpoint, rerun and replayed like the others (Appendix Q.6), gains $+0.1703$ of mean recall over the separately trained MOTIFS baseline: $+0.1460$ group-level (86%) and $+0.0243$, $[+0.0197, +0.0289]$, case-level. The MOTIFS baselines of two releases differ by $+0.0206$ in $mR@50$ with a case-level part of $+0.0007$, $[-0.0016, +0.0030]$ (last row of Tab. 1), so PE-Net's case-level part is not training noise; against the other baseline it is $+0.0236$. Its AUC is another matter: $+0.0050$, $[+0.0008, +0.0091]$, relation-weighted, but those baselines differ by -0.0060, $[-0.0099, -0.0011]$, and against the second PE-Net gains $+0.0110$, so we claim no AUC gain. Matched, its quadratic gain is almost all case-level ($+0.00609$ of $+0.00634$).

Beyond PredCls. In SGCls (58% of relations matched; Appendix Q.9) TDE raises $mR@50$ by $+0.0517$, 99% group-level, against the control's $+0.0518$ at $R@50 = 0.384$ (TDE's 0.263), and its relation-weighted AUC falls, -0.0128, $[-0.0171, -0.0074]$. In SGDet (65% matched) the control again gains more ($+0.0418$ against $+0.0310$).

6. The model the aggregates rank last

The predictors of Sec. 4 never see an image: they read the class pair and, at most, box coordinates from the annotation. On the same reconstructed split we add one that reads pixels — subject and object regions cropped from the image and encoded by a frozen CLIP ViT-B/32 [34], concatenated with the class embeddings MLP-CLASS uses. CLIP is not fine-tuned, so any case-level gain reflects what pretraining had already recovered. All three compared models are retrained on one fixed subsample of $100{,}000$ training relations, so differences isolate the features rather than the amount of data.

By every aggregate measure the CLIP model is the worst of the three: top-1 accuracy 0.6156 against 0.6467 for the geometry model, and a quadratic score 0.04642 lower. An evaluator reading either number would discard it. The split disagrees. Matched under the primary protocol, the CLIP model *gains* 0.00331 of case-level score against the geometry model, interval $[+0.00185, +0.00477]$ ($+0.00464$, $[+0.00311, +0.00618]$, by temperature alone; on raw outputs 0.00640, $[0.00513, 0.00767]$), while losing 0.04705 of group-level score under temperature matching. Its whole deficit is group-level (Fig. 2). Against the class-only model, pixels buy 0.01663 of case-level score on raw outputs where box geometry buys 0.01023, with well-separated intervals.

A larger case-level part is not the same as better ranking, and here the two disagree. With each model's prior matched within cells, the CLIP model still ranks the correct predicate first less often (top-1 accuracy 2.17 points lower, MRR 1.67 lower), and under the primary weighting the within-cell AUC favours geometry, -0.0067, $[-0.0113, -0.0015]$; by comparison the sign reverses, $+0.0085$, $[-0.0062, +0.0237]$. Over five training seeds (Appendix P) the CLIP model's temperature-matched advantage stays between $+0.00413$ and $+0.00474$, its accuracy below geometry's.

Under a label shift. Along a path to uniform predicate frequencies within each pair (Appendix Q.3), seven of the eight group-level parts change sign, the CLIP model beats geometry by $+0.02981$, and five of the six case-level parts that excluded zero keep their sign.

A debiasing gain that is case-level. On Waterbirds [37], with the background scene as the cell and linear probes on frozen CLIP features (Appendix Q.10), group reweighting against ERM raises worst-group accuracy from 0.643 to 0.854 and the quadratic score by $+0.0637$: a group-level loss of -0.0277 and a case-level gain of $+0.0914$, $[+0.0849, +0.0979]$, with the within-scene AUC up by $+0.0031$, $[+0.0008, +0.0056]$. Here debiasing buys what TDE does not, better discrimination inside each group. The supplement also splits CIFAR-100-LT and long-tailed text, where the fine label cannot be the cell: it forces the case-level part to zero.

7. Limitations and conclusion

Limitations. *The split is conditional on the grouping*: equally fine partitions can disagree (Appendix K), so a benchmark should declare ϕ before a leaderboard opens. *It credits any within-cell signal*, shortcuts included (Appendix I). *It moves with recalibration, the operating point and the label mix*; the AUC is invariant to the first two but not additive. *Its interval needs cells that are not tiny*: with two cases per cell it covers 81.5% (Appendix J). *It needs scores for every label*. And *each model is one released run*: intervals cover test sampling, not training.

Conclusion. Scoring each prediction against another case's label from the same group separates the two achievements mean recall sums. On three released prior-correcting PredCls checkpoints the gains are almost entirely group-level, and no weighting of TDE's within-pair AUC shows a gain; PE-Net, which changes features instead, has a case-level part well above the spread between separately trained baselines, though no AUC gain we can separate from it. We propose reporting, beside mean recall, the declared grouping, both parts with intervals, the within-group AUC and, for a thresholded metric, the operating-point control.

LLM Usage Statement

Large language models (Claude Opus, Sonnet and Haiku, Anthropic, used through the Claude Code agent) supported the writing and wrote code. For the writing, they

drafted and edited text, tables and figure layouts of the paper and the supplementary material, helped search and check the related literature and references, and produced simulated reviews used as feedback. For the code, they wrote and ran the analysis library, the replay of the official evaluator, the pipelines that re-ran the released checkpoints, the audit scripts and the script that checks every number in the manuscript against the stored results. The author reviewed and edited all text and code, checked the results, and takes full responsibility for the content.

References

[1] Aishwarya Agrawal, Dhruv Batra, Devi Parikh, and Aniruddha Kembhavi. Don't just assume; look and answer: Overcoming priors for visual question answering. In *CVPR*, pages 4971–4980, 2018. 3

[2] Jochen Bröcker. Reliability, sufficiency, and the decomposition of proper scores. *Quarterly Journal of the Royal Meteorological Society*, 135(643):1512–1519, 2009. 3

[3] Kaidi Cao, Colin Wei, Adrien Gaidon, Nikos Arechiga, and Tengyu Ma. Learning imbalanced datasets with label-distribution-aware margin loss. In *NeurIPS*, 2019. 3, 11, 20

[4] Chao Chen, Yibing Zhan, Baosheng Yu, Liu Liu, Yong Luo, and Bo Du. Resistance training using prior bias: Toward unbiased scene graph generation. In *AAAI*, pages 212–220, 2022. 3

[5] Tianshui Chen, Weihao Yu, Riquan Chen, and Liang Lin. Knowledge-embedded routing network for scene graph generation. In *CVPR*, pages 6156–6164, 2019. 2

[6] Meng-Jiun Chiou, Henghui Ding, Hanshu Yan, Changhu Wang, Roger Zimmermann, and Jiashi Feng. Recovering the unbiased scene graphs from the biased ones. In *ACM International Conference on Multimedia*, pages 1581–1590, 2021. 2, 7, 35

[7] Carlos Cinelli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B*, 82(1):39–67, 2020. 4, 14

[8] Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-balanced loss based on effective number of samples. In *CVPR*, pages 9268–9277, 2019. 11, 20

[9] Morris H. DeGroot and Stephen E. Fienberg. The comparison and evaluation of forecasters. *Journal of the Royal Statistical Society: Series D (The Statistician)*, 32(1-2):12–22, 1983. 3

[10] Elizabeth R. DeLong, David M. DeLong, and Daniel L. Clarke-Pearson. Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach. *Biometrics*, 44(3):837–845, 1988. 3, 14

[11] Alakh Desai, Tz-Ying Wu, Subarna Tripathi, and Nuno Vasconcelos. Learning of visual relations: The devil is in the tails. In *ICCV*, pages 15384–15393, 2021. 2

[12] Francis X Diebold and Roberto S Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995. 3, 4

[13] Xingning Dong, Tian Gan, Xuemeng Song, Jianlong Wu, Yuan Cheng, and Liqiang Nie. Stacked hybrid-attention and group collaborative learning for unbiased scene graph generation. In *CVPR*, pages 19405–19414, 2022. 2

[14] Aaron Fisher, Cynthia Rudin, and Francesca Dominici. All models are wrong, but many are useful: Learning a variable's importance by studying an entire class of prediction models simultaneously. *Journal of Machine Learning Research*, 20 (177):1–81, 2019. 3, 2

[15] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020. 3

[16] Raffaella Giacomini and Halbert White. Tests of conditional predictive ability. *Econometrica*, 74(6):1545–1578, 2006. 3

[17] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *CVPR*, pages 770–778, 2016. 20

[18] Wassily Hoeffding. A class of statistics with asymptotically normal distribution. *The Annals of Mathematical Statistics*, 19(3):293–325, 1948. 3, 4, 7

[19] Deunsol Jung, Sanghyun Kim, Won Hwa Kim, and Minsu Cho. Devil's on the edges: Selective quad attention for scene graph generation. In *CVPR*, pages 18664–18674, 2023. 2

[20] Corentin Kervadec, Grigory Antipov, Moez Baccouche, and Christian Wolf. Roses are red, violets are blue... but should vqa expect them to? In *CVPR*, pages 2775–2784, 2021. 2

[21] Ranjay Krishna, Yuke Zhu, Oliver Groth, Justin Johnson, Kenji Hata, Joshua Kravitz, Stephanie Chen, Yannis Kalantidis, Li-Jia Li, David A Shamma, et al. Visual genome: Connecting language and vision using crowdsourced dense image annotations. *International Journal of Computer Vision*, 123(1):32–73, 2017. 2, 15

[22] Ken Lang. Newsweeder: Learning to filter netnews. In *ICML*, pages 331–339, 1995. 22

[23] Lin Li, Long Chen, Yifeng Huang, Zhimeng Zhang, Songyang Zhang, and Jun Xiao. The devil is in the labels: Noisy label correction for robust scene graph generation. In *CVPR*, pages 18847–18856, 2022. 2

[24] Rongjie Li, Songyang Zhang, Bo Wan, and Xuming He. Bipartite graph network with adaptive message passing for unbiased scene graph generation. In *CVPR*, pages 11104–11114, 2021. 2, 37

[25] Rongjie Li, Songyang Zhang, Dahua Lin, Kai Chen, and Xuming He. From pixels to graphs: Open-vocabulary scene graph generation with vision-language models. In *CVPR*, 2024. 3

[26] Xingchen Li, Long Chen, Jian Shao, Shaoning Xiao, Songyang Zhang, and Jun Xiao. Rethinking the evaluation of unbiased scene graph generation. In *BMVC*, 2022. 2, 3

[27] Zachary C Lipton, Yu-Xiang Wang, and Alexander Smola. Detecting and correcting for label shift with black box predictors. In *ICML*, pages 3122–3130, 2018. 4

[28] Li Liu, Shuzhou Sun, Shuaifeng Zhi, Fan Shi, Zhen Liu, Janne Heikkilä, and Yongxiang Liu. A causal adjustment module for debiasing scene graph generation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(5): 4024–4043, 2025. 5

[29] Cewu Lu, Ranjay Krishna, Michael Bernstein, and Li Fei-Fei. Visual relationship detection with language priors. In *ECCV*, pages 852–869, 2016. 2, 37

[30] Xinyu Lyu, Lianli Gao, Yuyu Guo, Zhou Zhao, Hao Huang, Heng Tao Shen, and Jingkuan Song. Fine-grained predicates learning for scene graph generation. In *CVPR*, pages 19467–19475, 2022. 2

[31] Aditya Krishna Menon, Sadeep Jayasumana, Ankit Singh Rawat, Himanshu Jain, Andreas Veit, and Sanjiv Kumar. Long-tail learning via logit adjustment. In *ICLR*, 2021. 2, 3, 4, 5, 21

[32] Allan H. Murphy. A new vector partition of the probability score. *Journal of Applied Meteorology*, 12(4):595–600, 1973. 3

[33] Alexandre Perez-Lebel, Marine Le Morvan, and Gaël Varoquaux. Beyond calibration: Estimating the grouping loss of modern neural networks. In *International Conference on Learning Representations*, 2023. 3

[34] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *ICML*, pages 8748–8763, 2021. 8, 18, 39

[35] Paul R. Rosenbaum. *Observational Studies*. Springer, New York, 2 edition, 2002. 5, 14

[36] Marco Saerens, Patrice Latinne, and Christine Decaestecker. Adjusting the outputs of a classifier to new a priori probabilities: A simple procedure. *Neural Computation*, 14(1):21–41, 2002. 4

[37] Shiori Sagawa, Pang Wei Koh, Tatsunori B. Hashimoto, and Percy Liang. Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. In *International Conference on Learning Representations*, 2020. 8, 38

[38] Robert J. Serfling. *Approximation Theorems of Mathematical Statistics*. John Wiley & Sons, New York, 1980. 3, 9, 12

[39] Robik Shrestha, Kushal Kafle, and Christopher Kanan. A negative case analysis of visual grounding methods for VQA. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8172–8181, 2020. 3

[40] Stefan Siegert. Variance estimation for brier score decomposition. *Quarterly Journal of the Royal Meteorological Society*, 140(682):1771–1777, 2014. 3

[41] Carolin Strobl, Anne-Laure Boulesteix, Thomas Kneib, Thomas Augustin, and Achim Zeileis. Conditional variable importance for random forests. *BMC Bioinformatics*, 9:307, 2008. 3, 2

[42] Gopika Sudhakaran, Devendra Singh Dhami, Kristian Kersting, and Stefan Roth. Vision relation transformer for unbiased scene graph generation. In *ICCV*, pages 21882–21893, 2023. 2

[43] Kaihua Tang, Hanwang Zhang, Baoyuan Wu, Wenhan Luo, and Wei Liu. Learning to compose dynamic tree structures for visual contexts. In *CVPR*, pages 6612–6621, 2019. 2

[44] Kaihua Tang, Yulei Niu, Jianqiang Huang, Jiaxin Shi, and Hanwang Zhang. Unbiased scene graph generation from biased training. In *CVPR*, pages 3716–3725, 2020. 2, 5

[45] Damien Teney, Ehsan Abbasnejad, Kushal Kafle, Robik Shrestha, Christopher Kanan, and Anton van den Hengel. On the value of out-of-distribution testing: An example of Goodhart's law. In *Advances in Neural Information Processing Systems*, 2020. 3

[46] Yao Teng and Limin Wang. Structured sparse R-CNN for direct scene graph generation. In *CVPR*, pages 19415–19424, 2022. 2, 3

[47] Aad W. van der Vaart. *Asymptotic Statistics*. Cambridge University Press, Cambridge, UK, 1998. 9, 12

[48] Danfei Xu, Yuke Zhu, Christopher B Choy, and Li Fei-Fei. Scene graph generation by iterative message passing. In *CVPR*, pages 3097–3106, 2017. 2, 15

[49] Shaotian Yan, Chen Shen, Zhongming Jin, Jianqiang Huang, Rongxin Jiang, Yaowu Chen, and Xian-Sheng Hua. PCPL: Predicate-correlation perception learning for unbiased scene graph generation. In *ACM International Conference on Multimedia*, pages 265–273, 2020. 2

[50] J. Frank Yates. External correspondence: Decompositions of the mean probability score. *Organizational Behavior and Human Performance*, 30(1):132–156, 1982. 2, 3, 15

[51] Jing Yu, Yuan Chai, Yujing Wang, Yue Hu, and Qi Wu. CogTree: Cognition tree loss for unbiased scene graph generation. In *IJCAI*, pages 1274–1280, 2021. 2

[52] Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *International Conference on Learning Representations*, 2023. 40

[53] Rowan Zellers, Mark Yatskar, Sam Thomson, and Yejin Choi. Neural motifs: Scene graph parsing with global context. In *CVPR*, pages 5831–5840, 2018. 2

[54] Ao Zhang, Yuan Yao, Qianyu Chen, Wei Ji, Zhiyuan Liu, Maosong Sun, and Tat-Seng Chua. Fine-grained scene graph generation with data transfer. In *Computer Vision – ECCV 2022*, pages 409–424, 2022. 2, 7, 35

[55] Chaofan Zheng, Xinyu Lyu, Lianli Gao, Bo Dai, and Jingkuan Song. Prototype-based embedding network for scene graph generation. In *CVPR*, pages 22783–22792, 2023. 2, 7, 36

[56] Wenhao Zhou, Faqiang Liu, Hao Zheng, and Rong Zhao. Mitigating data bias and ensuring reliable evaluation of ai models with shortcut hull learning. *Nature Communications*, 16:5513, 2025. 3