

Highlights

WAGER: An Exact Decomposition of Model-Pair Score Gains on Benchmarks with Strong Label Priors

- Exact split of a model-pair score gain into label-prior fit and case-level evidence
- Audits released scene-graph checkpoints: a 10-point mean-recall gain is prior-fit
- Runs in $O(NK)$ on cached predictions; no baseline model, fold or tuning parameter
- Cluster-robust intervals, a limit theorem, and a group-stratified permutation test
- A CLIP model losing on every aggregate metric is shown to discriminate cases better