

WAGER: Within-cell Antisymmetric Gain Evaluation of Reasoning

Quang-Vinh Dang¹

¹British University Vietnam, Hung Yen, Vietnam.

Contributing authors: vinh.dq4@buv.edu.vn;

Abstract

Vision benchmarks with strong annotation priors can reward a new model even when its gain does not come from better instance-specific visual prediction. Existing audits usually compare a model with a frequency baseline, which confounds two questions: how well each system fits the prior, and whether the *difference* between systems uses information beyond it. We introduce WAGER, **W**ithin-cell **A**ntisymmetric **G**ain **E**valuation of **R**easoning, a new pairwise attribution algorithm. For two frozen probabilistic models, WAGER forms their proper-score contrast, matches examples that share a declared prior feature, and crosses their labels. This yields an exact finite-sample identity: observed model gain equals a prior-transported gain plus an antisymmetric instance-alignment gain. The latter is an order-two U-statistic and, with the quadratic score, equals twice the within-prior-cell prediction–label covariance gained by the new model. It is computed in $O(NK)$ time, requires no fitted prior model or projection fold, and admits image-cluster-robust inference and a cell-stratified randomization test. Controlled experiments recover injected instance signal monotonically and obtain type-I error **0.030** at level **0.05**. On 229,605 Visual Genome PredCls relations, a class-only MLP improves the frequency baseline entirely through the prior-transported channel, whereas adding box geometry produces an instance-alignment gain of **0.01439** (95% CI [**0.01373**, **0.01504**]; randomization $p = \mathbf{0.002}$). WAGER therefore attributes a model-to-model improvement directly, without treating a separately estimated baseline as ground truth.

Keywords: Visual Reasoning, Dataset Bias, Counterfactual Evaluation, U-statistics, Randomization Tests, Scene Graph Generation

WAGER: Within-cell Antisymmetric Gain Evaluation of Reasoning




Fig. 1 WAGER audits the gain directly. Prediction contrasts from two frozen models are matched within a declared prior cell, evaluated under observed and crossed label assignments, and separated exactly into prior-transported and instance-alignment components.

1 Introduction

Computer-vision benchmarks often contain a shortcut that is both powerful and easy to state. In scene graph generation (SGG), the ordered object-class pair strongly predicts the relation label; in visual question answering (VQA), question form predicts the answer; in long-tail recognition, category frequency dominates aggregate performance. Models can therefore improve a headline score by fitting annotation regularities more precisely without improving the instance-specific association between an image and its label. This is not merely hypothetical: frequency baselines are competitive on Visual Genome [1], VQA systems fail under changed answer priors [2, 3], and standard metrics obscure tail behavior [4].

The usual audit trains or constructs a prior-only baseline and compares its score with a full model. That strategy is useful diagnostically, but it does not identify where the *gain between two submitted systems* came from. Baseline estimation error, smoothing, calibration, and model-family mismatch all enter the result. More fundamentally, separately decomposing each model and subtracting the decompositions is an indirect answer to a pairwise question.

We instead ask a counterfactual question at the level of the gain itself. Let f_0 be the old model, f_1 the claimed improvement, and $\phi(x)$ the declared prior feature. For two test examples i, j with the same ϕ , we retain both models’ prediction vectors but cross the observed labels. The crossed assignment preserves the prior cell and its label histogram while destroying which prediction vector belongs with which label. Comparing the observed and crossed assignments measures whether $f_1 - f_0$ improves this instance-level alignment. Figure 1 summarizes the idea.

This construction gives WAGER: *Within-cell Antisymmetric Gain Evaluation of Reasoning*. “Reasoning” is used operationally: predictive alignment with the correct instance beyond ϕ , not a claim of human-like cognition or causality. WAGER makes four contributions.

1. We define a direct model-pair estimand whose prior component is generated by within-cell label transport, not by fitting a frequency forecaster.

2. We prove an exact population and finite-sample decomposition. The residual is an antisymmetric U-statistic; under the quadratic score it has a covariance identity that makes its meaning transparent.
3. We derive an $O(NK)$ algorithm, a Hájek influence function for image-cluster-robust intervals, and a cell-stratified randomization test under a sharp exchangeability null.
4. We validate the estimator under controlled alternatives and on Visual Genome, where it distinguishes a class-only gain from gains created by adding spatial evidence.

The redesign is important in what it removes. WAGER does not estimate a model’s “self-prior,” does not require a projection/evaluation split, does not equate betting growth with information, and does not rely on an arbitrary threshold for a gain ratio. The primary quantity is the signed instance-alignment gain with uncertainty.

2 Related Work

2.1 Benchmark priors and shortcut learning

Dataset bias has long complicated comparisons between visual recognition systems [5, 6]. In SGG, the FREQ predictor estimates $p(\text{predicate} \mid \text{subject}, \text{object})$ and is surprisingly competitive [1]; mean recall and causal interventions were proposed to reduce this dependence [7–9], although metric pathologies remain [4]. VQA-CP changes language priors between train and test [3], but changing-prior splits can themselves be exploited [10]. These methods alter training, data, or metrics. WAGER instead audits a fixed pair of probabilistic predictors on a fixed test set.

2.2 Information and predictive probing

Usable-information analyses measure information accessible to a chosen predictor family [11–13]. Their central objects are model- or family-level quantities. WAGER is deliberately narrower and more operational: it attributes the proper-score difference between two concrete frozen models. We do not identify conditional mutual information and do not claim that a positive residual establishes causal or compositional reasoning.

2.3 Proper scores, randomization, and U-statistics

Strictly proper scoring rules reward honest predictive distributions [14]; WAGER uses their model-pair contrast. Permutation methods have been used to assess classifier performance and conditional independence [15, 16]. Conditional permutation tests typically estimate or assume a conditional distribution when confounders are continuous or high-dimensional. WAGER exploits a different setting: the benchmark supplies a discrete audit cell ϕ , and the target is not generic conditional independence but a specific model-to-model score gain. This permits exact within-cell transport without estimating a nuisance law. The antisymmetric residual is an order-two U-statistic in the sense of Hoeffding [17], enabling an efficient all-pairs estimator and influence-function

inference. Pairwise betting has also been used to test exchangeability sequentially [18]; WAGER is a fixed-sample gain decomposition and makes no sequential-validity claim. Its new element is the gain-contrast transport identity and the resulting attribution algorithm, rather than permutation or U-statistic theory by itself.

3 Method

3.1 Setup: audit a gain, not a model

Let $q_0(\cdot | x)$ and $q_1(\cdot | x)$ be the predictive distributions of an old and a new frozen model over K labels. Let $S(q, y)$ be a bounded proper score, oriented so that larger is better, and let $C = \phi(X)$ be a declared discrete prior feature. For SGG, C is the ordered subject-object class pair. Define the instance-specific *gain vector*

$$H_x(y) = S(q_1(\cdot | x), y) - S(q_0(\cdot | x), y), \quad y = 1, \dots, K. \quad (1)$$

Unlike the realized gain $H_X(Y)$, this vector can be evaluated at a counterfactual label without rerunning either model.

Our default is the quadratic score

$$S_B(q, y) = 2q_y - \|q\|_2^2, \quad (2)$$

which differs from negative Brier loss only by a label-only constant. It is strictly proper and bounded, avoids probability clipping, and uses the entire predictive distribution. A floored log score is a sensitivity analysis.

3.2 Within-cell counterfactual transport

For a prior cell c , define three population quantities:

$$\Delta T_c = \mathbb{E}[H_X(Y) | C = c], \quad (3)$$

$$\Delta P_c = \mathbb{E}_{H|c} \mathbb{E}_{Y'|c}[H_X(Y')], \quad (4)$$

$$\Delta R_c = \Delta T_c - \Delta P_c. \quad (5)$$

Here Y' is an independent label drawn from the same cell. Thus ΔT_c is the observed score gain, ΔP_c is the gain retained after breaking the instance-specific prediction-label assignment while preserving the prior cell, and ΔR_c is the alignment gain lost by that transport.

Theorem 1 (Transport and covariance identities). *For any integrable score contrast H ,*

$$\Delta T_c = \Delta P_c + \Delta R_c, \quad \Delta R_c = \sum_{y=1}^K \text{Cov}(H_X(y), \mathbb{1}\{Y = y\} | C = c). \quad (6)$$




Fig. 2 The antisymmetric pair kernel. For two examples in the same prior cell, WAGER subtracts the score under crossed labels from the score under the observed assignment. Cell membership and label counts are unchanged; only instance alignment is removed.

For the quadratic score, writing $\Delta q_y(X) = q_1(y | X) - q_0(y | X)$,

$$\Delta R_c = 2 \sum_{y=1}^K \text{Cov}(\Delta q_y(X), \mathbb{I}\{Y = y\} \mid C = c). \quad (7)$$

Consequently, if the gain vector is a function of c alone, then $\Delta R_c = 0$.

Equation (7) is the operational meaning of WAGER: the new model receives reasoning credit only to the extent that its probability change aligns more strongly with the label of the correct instance within the same prior cell. Cell-constant score improvements, including improved fitting of the cell’s label histogram, remain in ΔP . The result does not assert that every nonzero covariance is causal reasoning; unrecorded within-cell shortcuts can also contribute, so ϕ must be declared and stress-tested.

3.3 Exact finite-sample decomposition

Consider cell c with $n_c \geq 2$ examples. Write $H_i(y) = H_{x_i}(y)$. WAGER computes

$$\widehat{\Delta T}_c = \frac{1}{n_c} \sum_{i \in c} H_i(y_i), \quad (8)$$

$$\widehat{\Delta P}_c = \frac{1}{n_c(n_c - 1)} \sum_{i \neq j \in c} H_i(y_j), \quad (9)$$

$$\widehat{\Delta R}_c = \binom{n_c}{2}^{-1} \sum_{i < j \in c} A_{ij}, \quad (10)$$

where the antisymmetric kernel is

$$A_{ij} = \frac{1}{2}\{H_i(y_i) + H_j(y_j) - H_i(y_j) - H_j(y_i)\}. \quad (11)$$

Theorem 2 (Finite-sample WAGER identity). *For every cell with $n_c \geq 2$, without a distributional assumption,*

$$\widehat{\Delta T}_c = \widehat{\Delta P}_c + \widehat{\Delta R}_c. \quad (12)$$

Moreover, under exchangeable sampling within the cell, $\widehat{\Delta R}_c$ is an unbiased order-two U -statistic for ΔR_c .

The dataset-level quantities weight cells by their number of identified examples: $\widehat{\Delta Z} = N_*^{-1} \sum_{c: n_c \geq 2} n_c \widehat{\Delta Z}_c$ for $Z \in \{T, P, R\}$. Singleton cells cannot identify a within-cell counterfactual; WAGER excludes and reports them instead of inventing a smoothed estimate. If $\widehat{\Delta T} > 0$, the descriptive alignment share is $\widehat{\rho} = \widehat{\Delta R} / \widehat{\Delta T}$. It may exceed one when instance alignment improves while prior fitting deteriorates, and it is undefined for a nonpositive total gain. We therefore report $\widehat{\Delta R}$ as the primary result, not a thresholded ratio.

3.4 Linear-time computation

Naively evaluating Eq. (9) is quadratic in cell size. Let m_{cy} be the count of label y in cell c . For example $i \in c$, its transported score is

$$P_i = \frac{\sum_{y=1}^K m_{cy} H_i(y) - H_i(y_i)}{n_c - 1}. \quad (13)$$

Averaging P_i yields Eq. (9); averaging $H_i(y_i) - P_i$ yields Eq. (10). Thus the entire estimator costs $O(NK)$ time and $O(NK)$ storage for cached prediction vectors. No model fitting, projection fold, hierarchy, or smoothing hyperparameter appears.

3.5 Inference

For point estimation we use all unordered within-cell pairs. For uncertainty, define $\bar{A}_i = (n_c - 1)^{-1} \sum_{j \neq i} A_{ij}$ and $\widehat{\Delta R}_c = n_c^{-1} \sum_i \bar{A}_i$. The estimated Hájek influence contribution for random cell membership is

$$\widehat{\psi}_i = 2\bar{A}_i - \widehat{\Delta R}_c - \widehat{\Delta R}. \quad (14)$$

We sum $\widehat{\psi}_i$ within image and use the resulting one-way sandwich variance, so multiple relations from the same image are not treated as independent. Total-gain influence is $H_i(y_i) - \widehat{\Delta T}$; prior-gain and share intervals follow from the exact identity and delta method.

As a complementary finite-group diagnostic, WAGER applies an independently sampled cyclic shift to labels inside every cell. Each shift preserves cell size and label

Algorithm 1 WAGER gain decomposition

Require: frozen q_1, q_0 ; labels y ; prior cells ϕ ; proper score S

Ensure: $\widehat{\Delta T}, \widehat{\Delta P}, \widehat{\Delta R}$ and uncertainty

```
1:  $H_i(k) \leftarrow S(q_{1i}, k) - S(q_{0i}, k)$  for all  $i, k$ 
2: for each cell  $c$  with  $n_c \geq 2$  do
3:   count labels  $m_{c1}, \dots, m_{cK}$ 
4:   for each  $i \in c$  do
5:      $P_i \leftarrow (\sum_k m_{ck} H_i(k) - H_i(y_i)) / (n_c - 1)$ 
6:      $R_i \leftarrow H_i(y_i) - P_i$ 
7:   end for
8: end for
9:  $\widehat{\Delta T} \leftarrow \text{mean } H_i(y_i)$ ;  $\widehat{\Delta P} \leftarrow \text{mean } P_i$ ;  $\widehat{\Delta R} \leftarrow \text{mean } R_i$ 
10: compute image-clustered Hájek interval and cell-shift randomization  $p$ 
```

histogram. Recomputing $\widehat{\Delta R}$ under B shifts gives

$$p = \frac{1 + \sum_{b=1}^B \mathbb{1}\{\widehat{\Delta R}^{(b)} \geq \widehat{\Delta R}\}}{B + 1}. \quad (15)$$

This is exact under the sharp null that labels are exchangeable relative to the frozen model outputs within each cell. Because SGG relations share images, our primary real-data uncertainty statement is the image-cluster-robust interval; the randomization test is reported as a secondary relation-level diagnostic.

4 Experiments

We evaluate three questions: (i) does WAGER remain centered under the null and respond monotonically to known instance signal; (ii) does it distinguish class-prior fitting from spatial evidence on a real long-tailed benchmark; and (iii) are conclusions stable to score, cell granularity, and sparse-cell filtering? All confidence intervals are 95%.

4.1 Controlled validation

We generate 24 prior cells with five labels and draw each cell’s label distribution from a Dirichlet law. The old model is exactly the cell prior. The new model is $q_1 = (1 - \beta)q_0 + \beta q_{\text{oracle}}$, where q_{oracle} assigns 0.98 to the realized label. Forty repetitions are run at each $\beta \in \{0, .1, \dots, .8\}$. The estimated alignment gain grows linearly and monotonically with the injected signal (Figure 3, left).

For type-I evaluation, we add mean-zero perturbations to prior predictions and randomly reassign the perturbed prediction vectors within each cell, ensuring no systematic prediction–label alignment. Across 200 datasets, the 99-shift randomization test rejects in 0.030 of runs at level 0.05; the mean estimated alignment gain is -8.97×10^{-5} . The null p -values are approximately uniform and null estimates remain centered at zero (Figure 3, center and right).




Fig. 3 Controlled validation. WAGER recovers injected within-cell signal monotonically (mean $\pm$ one standard deviation), produces approximately uniform null randomization p -values, and controls type-I error in 200 null runs.

Table 1 Frozen-model top-1 accuracy on VG150 PredCls. Similar accuracies do not reveal whether the difference is prior-recoverable.

Model	Accuracy
FREQ	0.6564
FREQ+OVERLAP	0.6545
MLP-CLASS	0.6554
MLP-SPATIAL	0.6639
MLP-SPATIAL+	0.6674

4.2 Visual Genome setup

We use Visual Genome [19] in the standard VG150 PredCls regime: ground-truth boxes and object labels are supplied, and the model predicts one of 50 predicate labels for each ordered object pair. The training/test split contains 532,987/229,605 relations from 65,228/27,955 images. The prior cell is $\phi = (\text{subject class}, \text{object class})$. Of 8,614 test cells, 6,346 contain at least two relations. WAGER identifies 227,337 test relations (99.0%); the 2,268 singleton cells are reported and excluded.

Five deliberately controlled predictors expose different information. FREQ is the training-count distribution conditioned on the class pair. FREQ+OVERLAP additionally uses a binary box-overlap indicator. MLP-CLASS consumes fixed subject/object class embeddings and is therefore a function of ϕ alone. MLP-SPATIAL adds 12 normalized box-geometry features (relative center, scale, aspect, intersection, containment, and IoU); MLP-SPATIAL+ increases its capacity. All models are trained once on the training split. WAGER consumes only their cached test distributions. We use the quadratic score, image-clustered intervals, and 499 within-cell cyclic randomizations.

Table 2 WAGER decomposition on 227,337 identified VG150 PredCls relations. Gains use the quadratic score. CI is image-cluster robust; p is the 499-shift relation-level randomization diagnostic. “–” denotes an undefined share because total gain is negative.

New model	Old model	$\widehat{\Delta T}$	$\widehat{\Delta P}$	$\widehat{\Delta R}$ (95% CI)	$\hat{\rho}$	p
FREQ+OVERLAP	FREQ	−0.01229	−0.01430	0.00202 [0.00168, 0.00235]	–	.002
MLP-CLASS	FREQ	0.03386	0.03386	0.00000 [0.00000, 0.00000]	0.00	.724
MLP-SPATIAL	FREQ	0.04270	0.02832	0.01439 [0.01373, 0.01504]	0.34	.002
MLP-SPATIAL+	FREQ	0.04644	0.02885	0.01760 [0.01686, 0.01834]	0.38	.002
MLP-SPATIAL	MLP-CLASS	0.00885	−0.00554	0.01439 [0.01373, 0.01504]	1.63	.002
MLP-SPATIAL+	MLP-SPATIAL	0.00374	0.00053	0.00321 [0.00289, 0.00353]	0.86	.002




Fig. 4 Direct model-pair attribution. Left: exact decomposition of each positive gain into prior-transported and instance-alignment components (the negative prior term for SPATIAL-CLASS extends left of zero). Right: primary alignment estimates with image-cluster-robust intervals.

4.3 Direct gain attribution

Table 2 and Figure 4 give the central result. MLP-CLASS improves the quadratic score over FREQ by 0.03386, but its gain vector is constant inside every class-pair cell. WAGER therefore assigns the entire improvement to prior transport and exactly zero to instance alignment. This is an algebraic consequence of the estimator, not a learned baseline passing a sanity check.

MLP-SPATIAL versus FREQ gains 0.04270 in total: 0.02832 is retained after label transport and 0.01439 is instance alignment (CI [0.01373, 0.01504]). The direct comparison MLP-SPATIAL versus MLP-CLASS is more revealing. Its total gain is only 0.00885, because prior fitting deteriorates by 0.00554, while spatial alignment improves by 0.01439. Thus the alignment share is 1.63: the useful instance-specific change is partially hidden by a worse prior component. A ratio constrained to [0, 1] or thresholded at 0.5 would misdescribe this tradeoff.

MLP-SPATIAL+ adds 0.00374 over MLP-SPATIAL, of which 0.00321 is alignment (CI [0.00289, 0.00353]). FREQ+OVERLAP illustrates the opposite caution: overlap creates a small positive alignment term, but its negative prior term is larger, so the model is worse overall. WAGER does not label any use of image-derived information as an improvement; it reports both channels and their sign.

4.4 Sensitivity analyses

For MLP-SPATIAL versus MLP-CLASS, the conclusion survives every planned sensitivity check. Changing from the quadratic to a probability-floored log score changes scale (0.01439 to 0.06536) but not sign or significance. Coarsening ϕ from the class pair to subject class increases the credited alignment to 0.02181, as expected: the coarser audit removes less prior structure and is less conservative. Raising the minimum cell size from 2 to 5, 10, and 20 yields 0.01395, 0.01330, and 0.01225; all cluster-robust lower limits remain positive. Unlike the rejected projection-based formulation, there is no smoothing strength, split fraction, clipping constant, betting order, or gain-ratio decision threshold to tune.

5 Discussion and Limitations

WAGER changes the unit of analysis from a model to a model-pair gain. This matters scientifically. A class-only network may be a better prior estimator than FREQ, but that fact does not become evidence of instance-specific visual use. Conversely, a model can improve visual alignment while losing enough prior score to obscure the improvement in its aggregate metric. The signed identity $\Delta T = \Delta P + \Delta R$ exposes both cases without a separately fitted reference forecaster.

The interpretation remains conditional on the declared audit feature ϕ . A positive ΔR means the new model aligns its probability changes with the correct example better than the old model among examples sharing ϕ . It does not prove causal reasoning, compositional understanding, or robustness. If an unrecorded shortcut varies inside a cell, WAGER will count its predictive use in ΔR . Applications should therefore choose ϕ before looking at results, justify it from the benchmark’s data generation process, and report finer and coarser alternatives. The observed increase under subject-only cells illustrates why the finest defensible prior feature is the conservative primary choice.

Exact transport also has a positivity requirement: a cell needs at least two examples. WAGER reports the identified fraction and makes no claim for singletons. Very fine ϕ can therefore improve confound control while reducing coverage. Approximate matching could extend the method to continuous priors, but would replace the exact identity’s clear intervention with modeling assumptions; we leave that extension open.

The image-clustered confidence interval is asymptotic. The cyclic randomization test is finite-group exact only under its sharper within-cell exchangeability null and treats relations as the randomized units, so it is secondary in the SGG analysis. A future image-level randomization design that preserves several relation-cell margins simultaneously would strengthen finite-sample inference for structured benchmarks.

Finally, WAGER evaluates probability vectors, not only top-1 decisions. This is a feature—the method can detect useful redistribution below the argmax—but it requires released logits or probabilities. The quadratic score is bounded and strictly proper; other proper scores answer scale-dependent variants of the same question and should be declared in advance.

6 Broader applicability

The algorithm needs only four aligned arrays: two models’ probability vectors, labels, and prior-cell identifiers. For VQA, ϕ can be a predeclared question template; for long-tail recognition, it can encode collection source and frequency tier; for multimodal benchmarks, it can encode answer format or prompt family. In each case the scientific claim must remain precise: WAGER measures new prediction–label alignment within those cells. Its value is not that one choice of ϕ defines reasoning universally, but that every attribution becomes an explicit, reproducible counterfactual tied to a declared nuisance structure.

7 Conclusion

We introduced WAGER, Within-cell Antisymmetric Gain Evaluation of Reasoning, a direct algorithm for attributing the score improvement between two vision models. By crossing labels among examples with the same prior feature, WAGER constructs the prior-only counterfactual from the evaluated data themselves and obtains an exact identity between total, prior-transported, and instance-alignment gains. The residual is an interpretable U-statistic, equals a conditional covariance gain under the quadratic score, and is computable in $O(NK)$ time. Controlled experiments support null calibration and signal recovery. On Visual Genome, the method assigns a class-only improvement entirely to the prior channel and identifies the added alignment produced by spatial geometry. This gain-specific mechanism, rather than a combination of a frequency baseline and a generic significance test, is the paper’s technical contribution.

Data Availability

All code and data used in this study are publicly available at <https://github.com/vinhqdang/WAGER>.

References

- [1] Zellers, R., Yatskar, M., Thomson, S., Choi, Y.: Neural motifs: Scene graph parsing with global context. In: CVPR, pp. 5831–5840 (2018)
- [2] Goyal, Y., Khot, T., Summers-Stay, D., Batra, D., Parikh, D.: Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In: CVPR, pp. 6904–6913 (2017)
- [3] Agrawal, A., Batra, D., Parikh, D., Kembhavi, A.: Don’t just assume; look and answer: Overcoming priors for visual question answering. In: CVPR, pp. 4971–4980 (2018)
- [4] Li, W., Zhang, H., Bai, Q., Zhao, G., Jiang, N., Yuan, X.: Rethinking the evaluation of unbiased scene graph generation. BMVC (2022)

- [5] Torralba, A., Efros, A.A.: Unbiased look at dataset bias. In: CVPR, pp. 1521–1528 (2011)
- [6] Geirhos, R., Jacobsen, J.-H., Michaelis, C., Zemel, R., Brendel, W., Bethge, M., Wichmann, F.A.: Shortcut learning in deep neural networks. *Nature Machine Intelligence* **2**(11), 665–673 (2020)
- [7] Chen, T., Yu, W., Chen, R., Lin, L.: Knowledge-embedded routing network for scene graph generation. In: CVPR, pp. 6163–6171 (2019)
- [8] Tang, K., Zhang, H., Wu, B., Luo, W., Liu, W.: Learning to compose dynamic tree structures for visual contexts. In: CVPR, pp. 6619–6628 (2019)
- [9] Tang, K., Niu, Y., Huang, J., Shi, J., Zhang, H.: Unbiased scene graph generation from biased training. In: CVPR, pp. 3716–3725 (2020)
- [10] Teney, D., Kafle, K., Shrestha, R., Abbasnejad, E., Kanan, C., Hengel, A.: On the value of out-of-distribution testing: An example of goodhart’s law. In: NeurIPS (2020)
- [11] Xu, Y., Zhao, S., Song, J., Stewart, R., Ermon, S.: A theory of usable information under computational constraints. In: ICLR (2020)
- [12] Ethayarajh, K., Choi, Y., Swayamdipta, S.: Understanding dataset difficulty with v-usable information. In: ICML, pp. 5988–6008 (2022). PMLR
- [13] Hewitt, J., Ethayarajh, K., Liang, P., Manning, C.D.: Conditional probing: measuring usable information beyond a baseline. In: EMNLP, pp. 1626–1639 (2021)
- [14] Gneiting, T., Raftery, A.E.: Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association* **102**(477), 359–378 (2007) <https://doi.org/10.1198/016214506000001437>
- [15] Ojala, M., Garriga, G.C.: Permutation tests for studying classifier performance. *Journal of Machine Learning Research* **11**, 1833–1863 (2010)
- [16] Berrett, T.B., Wang, Y., Barber, R.F., Samworth, R.J.: The conditional permutation test for independence while controlling for confounders. *Journal of the Royal Statistical Society: Series B* **82**(1), 175–197 (2020)
- [17] Hoeffding, W.: A class of statistics with asymptotically normal distribution. *The Annals of Mathematical Statistics* **19**(3), 293–325 (1948)
- [18] Saha, A., Ramdas, A.: Testing exchangeability by pairwise betting. In: Proceedings of the 27th International Conference on Artificial Intelligence and Statistics, vol. 238, pp. 4915–4923. PMLR, Valencia, Spain (2024)

- [19] Krishna, R., Zhu, Y., Groth, O., Johnson, J., Hata, K., Kravitz, J., Chen, S., Kalantidis, Y., Li, L.-J., Shamma, D.A., *et al.*: Visual genome: Connecting language and vision using crowdsourced dense image annotations. International Journal of Computer Vision **123**(1), 32–73 (2017)

Appendix A Proofs

A.1 Proof of Theorem 1

Condition throughout on $C = c$. Expanding over labels gives

$$\Delta T_c = \sum_y \mathbb{E}[H_X(y) \mathbb{I}\{Y = y\}], \quad \Delta P_c = \sum_y \mathbb{E}[H_X(y)] \Pr(Y = y).$$

Subtracting yields the covariance sum in Eq. (6); the first identity is the definition of ΔR_c . Under the quadratic score,

$$H_X(y) = 2\Delta q_y(X) - \{\|q_1(\cdot | X)\|_2^2 - \|q_0(\cdot | X)\|_2^2\}.$$

The bracketed term is independent of the label index. Its contribution summed over y is $\text{Cov}(B(X), \sum_y \mathbb{I}\{Y = y\}) = \text{Cov}(B(X), 1) = 0$. The remaining terms give Eq. (7). If $H_X(y) = H_c(y)$ almost surely, every covariance is zero. $\square$

A.2 Proof of Theorem 2

Summing the kernel in Eq. (11) over unordered pairs, every observed term $H_i(y_i)$ occurs in $n_c - 1$ pairs with coefficient $1/2$, while the two crossed terms together enumerate every ordered pair $H_i(y_j)$, $i \neq j$, with coefficient $1/2$. Since $\binom{n_c}{2} = n_c(n_c - 1)/2$,

$$\widehat{\Delta R}_c = \frac{1}{n_c} \sum_i H_i(y_i) - \frac{1}{n_c(n_c - 1)} \sum_{i \neq j} H_i(y_j) = \widehat{\Delta T}_c - \widehat{\Delta P}_c.$$

This is deterministic. Under exchangeable sampling, the average of the symmetric order-two kernel is the standard unbiased U-statistic for its population expectation [17]. $\square$

Appendix B Efficient influence computation

The implementation does not instantiate all pairs. In addition to Eq. (13), let $G_c(y) = \sum_{j \in c} H_j(y)$ and $O_c = \sum_{j \in c} H_j(y_j)$. The pair-kernel mean for observation i is

$$\bar{A}_i = \frac{1}{2} \left[H_i(y_i) + \frac{O_c - H_i(y_i)}{n_c - 1} - P_i - \frac{G_c(y_i) - H_i(y_i)}{n_c - 1} \right].$$

Therefore both the point estimate and Hájek projection require only label counts, column sums, and matrix–vector products inside each cell. For image clusters g , the implemented variance is

$$\widehat{\text{Var}}(\widehat{\Delta R}) = \frac{G}{G-1} \frac{1}{N_*^2} \sum_{g=1}^G \left(\sum_{i \in g} \widehat{\psi}_i \right)^2,$$

where G is the number of images and $\widehat{\psi}_i$ is Eq. (14). We use a normal critical value; with 27,955 image clusters, small- G corrections are immaterial.

Appendix C Randomization algorithm

For each cell, sort examples once and index them $0, \dots, n_c - 1$. Randomization b draws $o_c^{(b)}$ uniformly from this index set and maps label position r to $(r + o_c^{(b)}) \bmod n_c$. Independent offsets across cells define an element of the product of cyclic groups. The cell label histogram is unchanged, so $\sum_y m_{cy} H_i(y)$ in Eq. (13) is precomputed; only the assigned label lookup changes. One randomization is $O(N)$ after the $O(NK)$ setup. The plus-one Monte Carlo correction includes the observed assignment and prevents zero p -values.

Appendix D Implementation and reproducibility

The NumPy implementation is in `wager/antisymmetric.py`. The primary routine `decompose_gain` validates probability matrices, forms score contrasts, excludes singleton cells, computes all three components, verifies the identity numerically, and returns per-instance transported/alignment values plus intervals. The Visual Genome driver is `experiments/run_sgg_wager.py`; controlled validation and sensitivity checks are in `experiments/antisymmetric_simulation.py` and `experiments/antisymmetric_ablations.py`. All random seeds are fixed. Unit tests cover the decomposition identity, quadratic covariance identity, cell-constant null, singleton reporting, cluster inference, randomization power, and log-score path.